

# *Realtime* Time Series Anomaly Detection in Production



Shreya Khurana  
Data Scientist - Intuit

# Intro



## **Intuit Identity Analytics Team:**

- Derek Schwartz, Manager 3 Data and Analytics
- Aaron Walker, Principal Technical Data Analyst
- Jacob Langley, Staff Technical Data Analyst

## **Data Science Team:**

- April Liu, Manager 2 Data Science

My team

## **Data Scientist at Intuit Bay Area, California, USA**

I work on the anomaly detection capability that tracks authentication and business health metrics at Intuit.

Previously I was building deep learning models to improve domain name recommendations at GoDaddy.

I have a masters in statistics from UIUC.

I'm a Python enthusiast and enjoy sharing my learnings with the community

What I do

intuit

GoDaddy



# What This Talk Is About

*Realtime* Time Series

Anomaly Detection


in Production

# Time Series Data and Applications

- Marketing / sales
- Weather
- KPI monitoring
- Data health checks
- Predictive maintenance
- Any monitoring system

Time series data: data measuring the same thing over a period of time

**Exhibit some behavior like trend or seasonality that can be modeled statistically.**



Market Summary > NVIDIA Corp


138.25 USD

+91.48 (195.60%) ↑ past year


Closed: 29 Nov, 5:00 pm GMT-5 • Disclaimer

After hours 138.14 -0.11 (0.080%)

1D 5D 1M 6M YTD 1Y 5Y Max

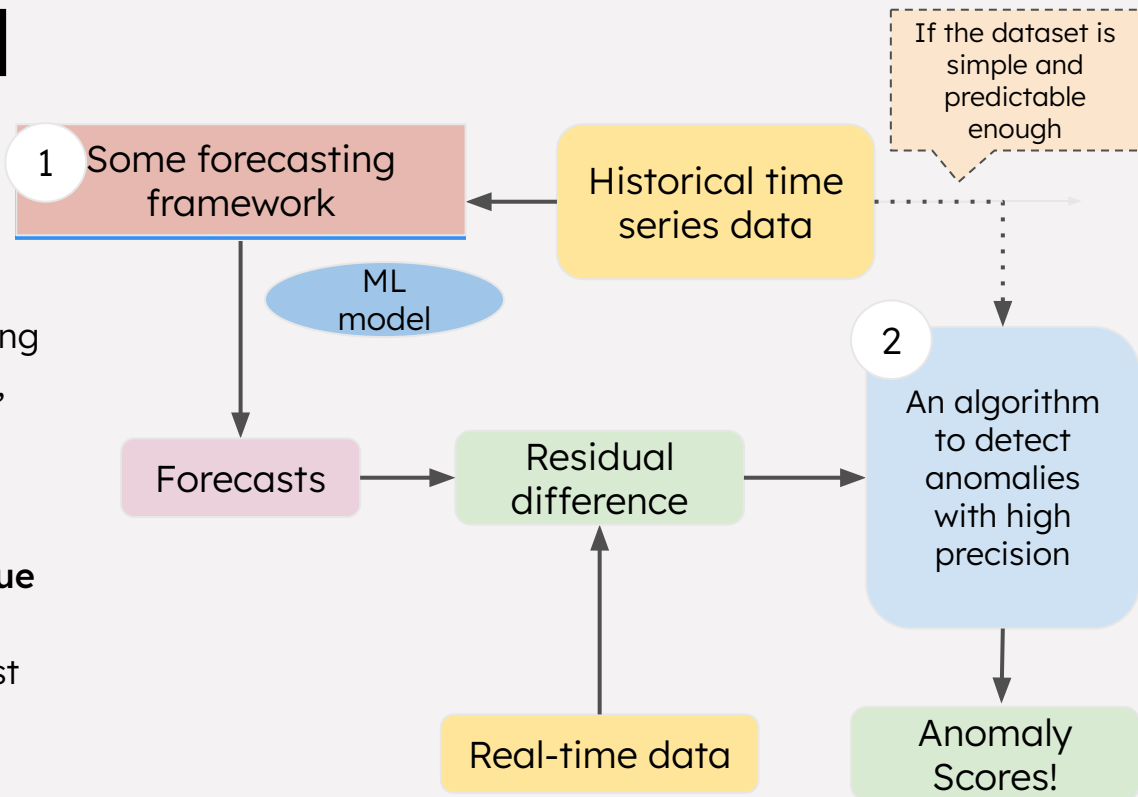


Monthly Cold and Flu Product Sales in 2020




# Anomaly Detection on Time Series Data

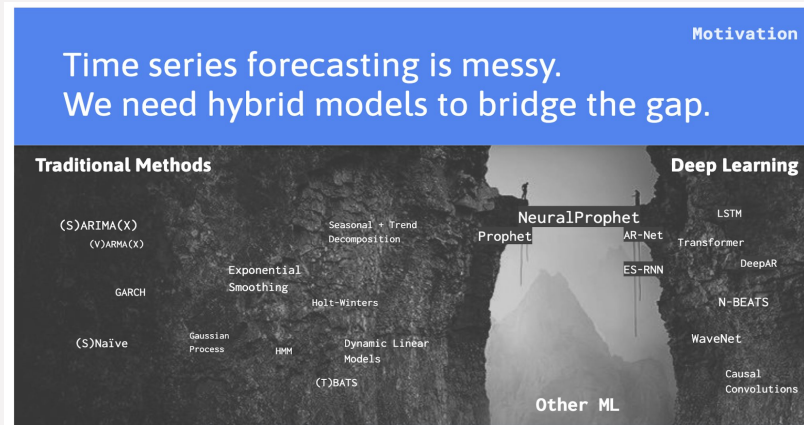
- Anomaly is when something is different that the expected or the usual
- Start simple with threshold based alerting
- If your time series is more complicated, then:
  - Assume your future data will behave *somewhat* similar to past
  - Forecast to get the **expected value**
  - Alert if the incoming data is different enough from the forecast
  - Need two core pieces



# NeuralProphet: A Time Series Modeling Framework



Classical time series models had a limitation with modeling long-range dependencies



NeuralProphet is a scalable framework that uses a neural network to model autoregression

# NeuralProphet Demo

## Model fitting and prediction

Component Decomposition

Modeling Custom Events

Tuning Model Parameters

```
from neuralprophet import NeuralProphet
```

```
m = NeuralProphet()
```

```
# Fit the model on the dataset (this might take a bit)
```

```
metrics = m.fit(df)
```

```
# Create a new dataframe reaching into the future for our forecast,
```

```
n_historic_predictions also shows historic data
```

```
df_future = m.make_future_dataframe(df, n_historic_predictions=True, periods=24*7)
```


```
# Predict
```

```
forecast = m.predict(df_future)
```

```
# Visualize the forecast
```

```
fig = m.plot(forecast)
```

```
fig.show()
```



# NeuralProphet Demo

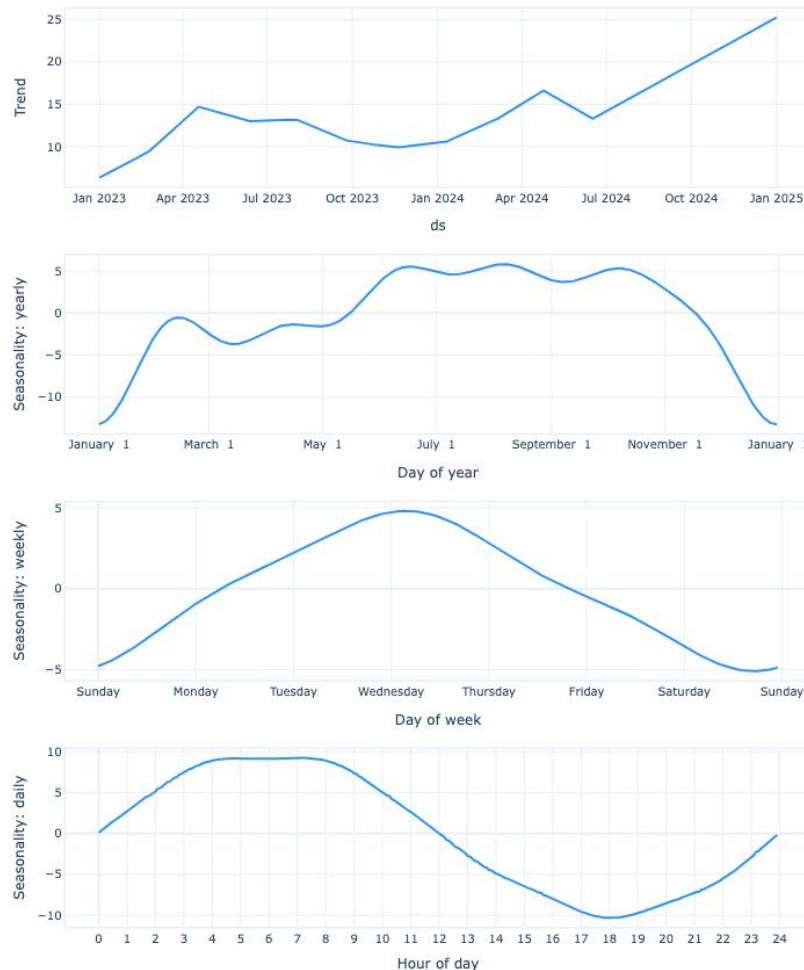
Model fitting and prediction

## Component Decomposition

Modeling Custom Events

Tuning Model Parameters

```
fig = m.plot_parameters()  
fig.show()
```

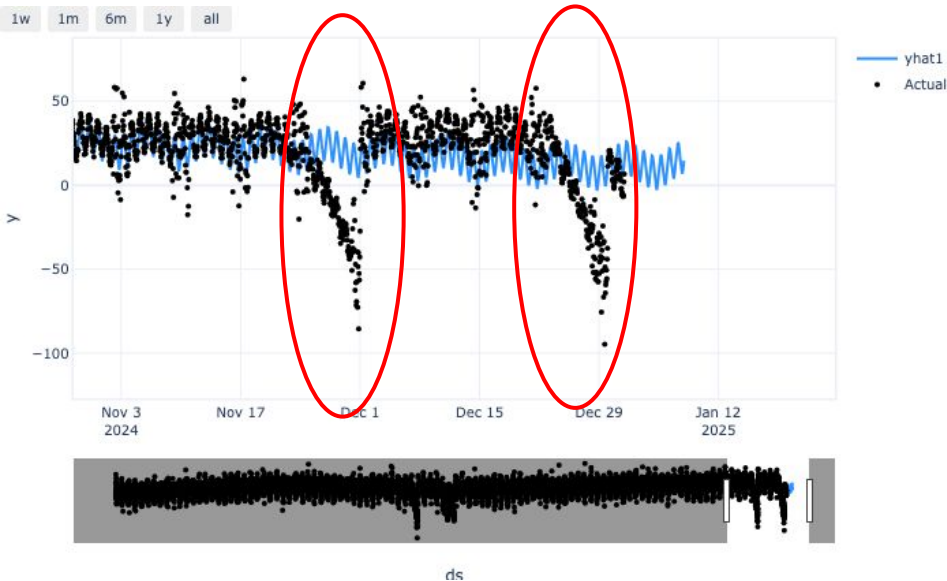




# NeuralProphet Demo

Model fitting and prediction  
Component Decomposition  
**Modeling Custom Events**  
Tuning Model Parameters

Before:



Observe the model doesn't fit very well during holiday periods by default.

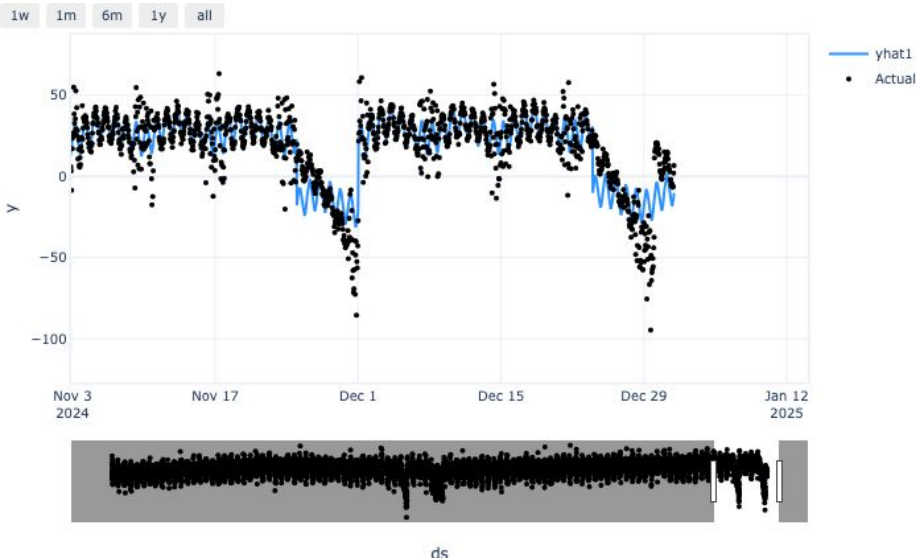
We need to add this signal to the model explicitly.

# NeuralProphet Demo

Model fitting and prediction  
Component Decomposition  
**Modeling Custom Events**  
Tuning Model Parameters

```
holiday_weeks = {  
    "Thanksgiving": ["2023-11-20/2023-11-26", "2024-11-25/2024-12-01"],  
    "Christmas": ["2023-12-23/2023-12-29", "2024-12-24/2024-12-30"],  
    "NewYear": ["2023-12-30/2024-01-05", "2024-12-30/2025-01-05"]  
}  
  
df_events = pd.DataFrame()  
for holiday, periods in holiday_weeks.items():  
    for period in periods:  
        start_date, end_date = (pd.to_datetime(p) for p in period.split('/'))  
        df_events = pd.concat([df_events, pd.DataFrame({"event": "holiday", "ds":  
pd.date_range(start_date, end_date, freq="h"))})]  
  
m.add_events("holiday")  
  
df_all = m.create_df_with_events(df, df_events)  
  
metrics = m.fit(df_all)  
forecast = m.predict(df_all)
```

After:




# NeuralProphet Demo


Model fitting and prediction  
Component Decomposition  
Modeling Custom Events  
**Tuning Model Parameters**

Smart defaults for parameters and hyperparameters,  
but can be manually overridden


```
fig = m.plot_parameters(components=["trend"])  
fig.show()
```



```
m = NeuralProphet(n_changepoints=2)  
m.fit(df)  
fig = m.plot_parameters(components=["trend"])  
fig.show()
```



# Anomaly Detection



# Z-scores

Use z-scores to determine the extent of an anomaly i.e. how far a data point is from the average.

---

## Mean based

---

**Mean** of the distribution from some period

$$z = \frac{X - \mu}{\sigma}$$

---


Measure spread of the distribution through **standard deviation**

# Lookback Period

Period which you “look back upon”  
to find how the non-anomalous  
distribution should look like

Calculate stats like averages and  
spread of the distribution


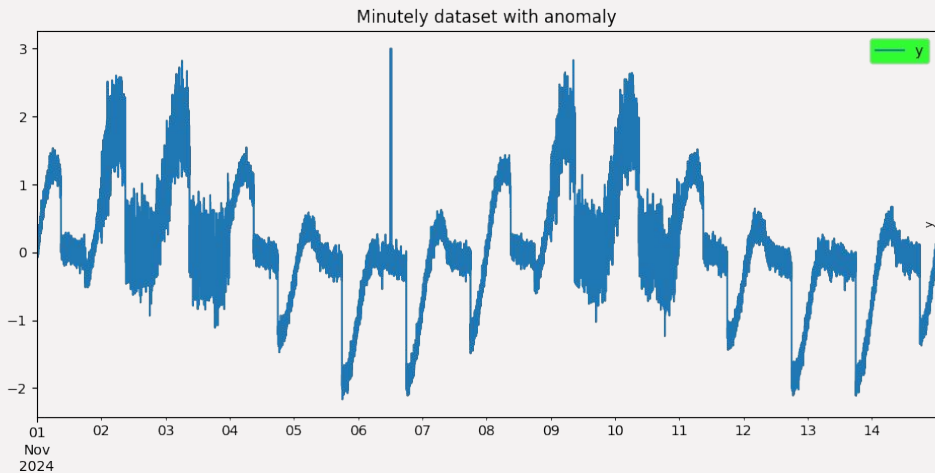
- What makes most sense?
- Experiment experiment experiment!
- Try multiple windows and domain knowledge
- Latency!
- Example: think about seasonality and hour of day and day of week



# Lookback Period

## Using the entire history

```
mean = minutely_df['y'].mean()  
std = minutely_df['y'].std()  
minutely_df['z_score'] = (minutely_df['y'] - mean) / std
```



# Lookback Period


Using last 10 minutes

```
minutely_df['rolling_mean'] =  
minutely_df['y'].rolling(window=10).mean()
```


```
minutely_df['rolling_std'] =  
minutely_df['y'].rolling(window=10).std()
```

```
minutely_df['z_score'] = (minutely_df['y'] -  
minutely_df['rolling_mean']) / minutely_df['rolling_std']
```

Minutely dataset with anomaly



Minutely dataset with anomaly score





# Lookback Period


## Partitioning weekdays and weekends

```
minutely_df['weekend'] = minutely_df['ds'].dt.dayofweek >= 5
```



```
weekend_mean = minutely_df[minutely_df['weekend']][['y']].mean()  
weekend_std = minutely_df[minutely_df['weekend']][['y']].std()  
weekday_mean = minutely_df[~minutely_df['weekend']][['y']].mean()  
weekday_std = minutely_df[~minutely_df['weekend']][['y']].std()
```

```
minutely_df['z_score'] = np.where(minutely_df['weekend'],  
                                   (minutely_df['y'] - weekend_mean) / weekend_std,  
                                   (minutely_df['y'] - weekday_mean) / weekday_std)
```

Minutely dataset with anomaly



Minutely dataset with anomaly score



# Z-scores

Use z-scores to determine how far a data point is from the average.

- An “average” can be defined in multiple ways
- We also need a way to measure spread of the data, to put this distance in context

---

## Mean based

---

**Mean** of the distribution from some period

Measure spread of the distribution through  
**standard deviation**

---

## Median based

---


**Median** of the distribution from some period

Measure spread of the distribution through  
**median absolute deviation**

# Z-scores

Use z-scores to determine how far a data point is from the average.

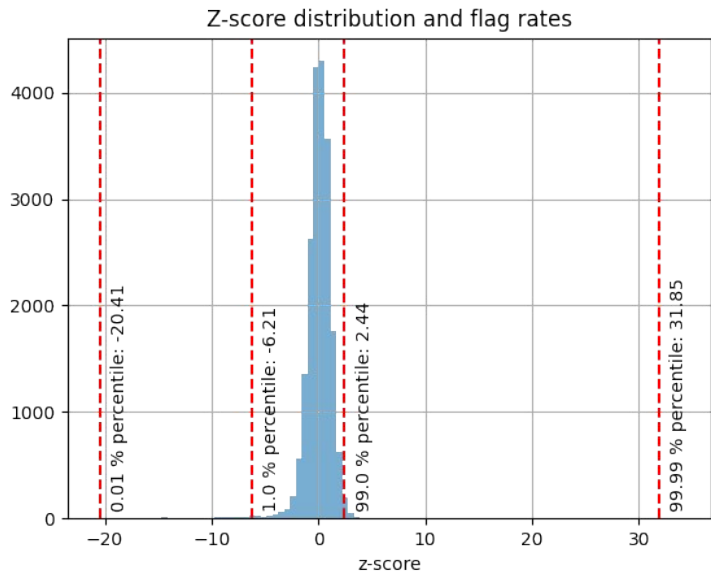
- An “average” can be defined in multiple ways
- We also need a way to measure spread of the data, to put this distance in context



# Tuning anomaly scores

How do we know if the data is “far enough” to call it anomalous?


- If you have ground truth of anomalies, amazing!
- If not:
  - Visual inspection
  - Flag a given percentile (say, 0.001% of the data)
  - Flag and run for some time
  - Measure precision/recall



# Tuning anomaly scores

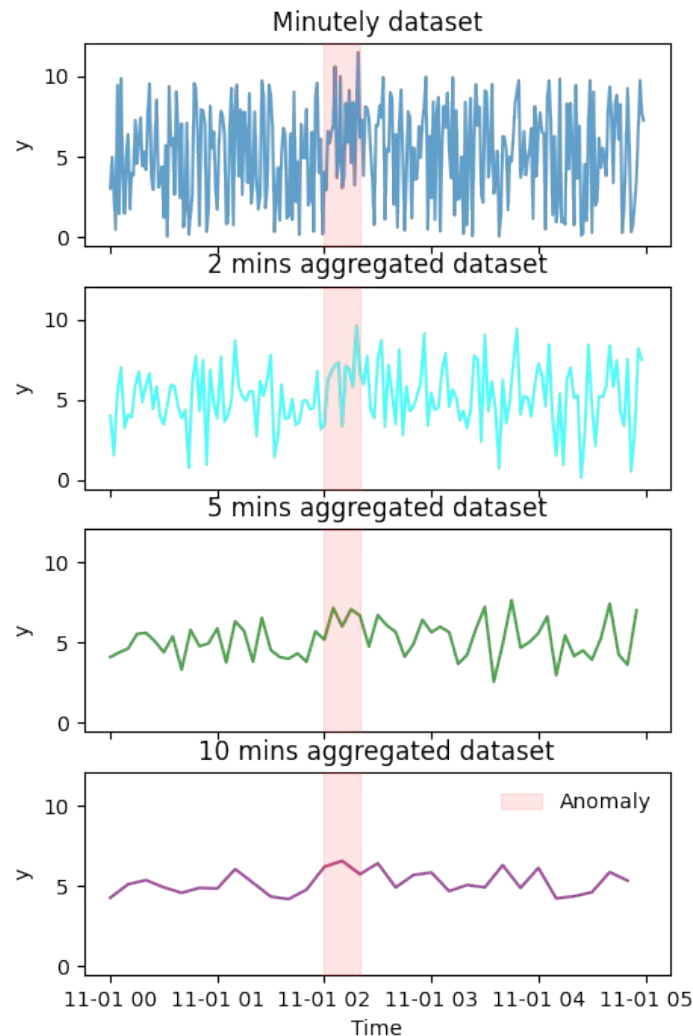
How do we set thresholds in case of a volatile time series and volatile z-scores?

- Smoothing of scores
  - Average anomaly scores over a period of time
- Accuracy ve time-to-detect latency



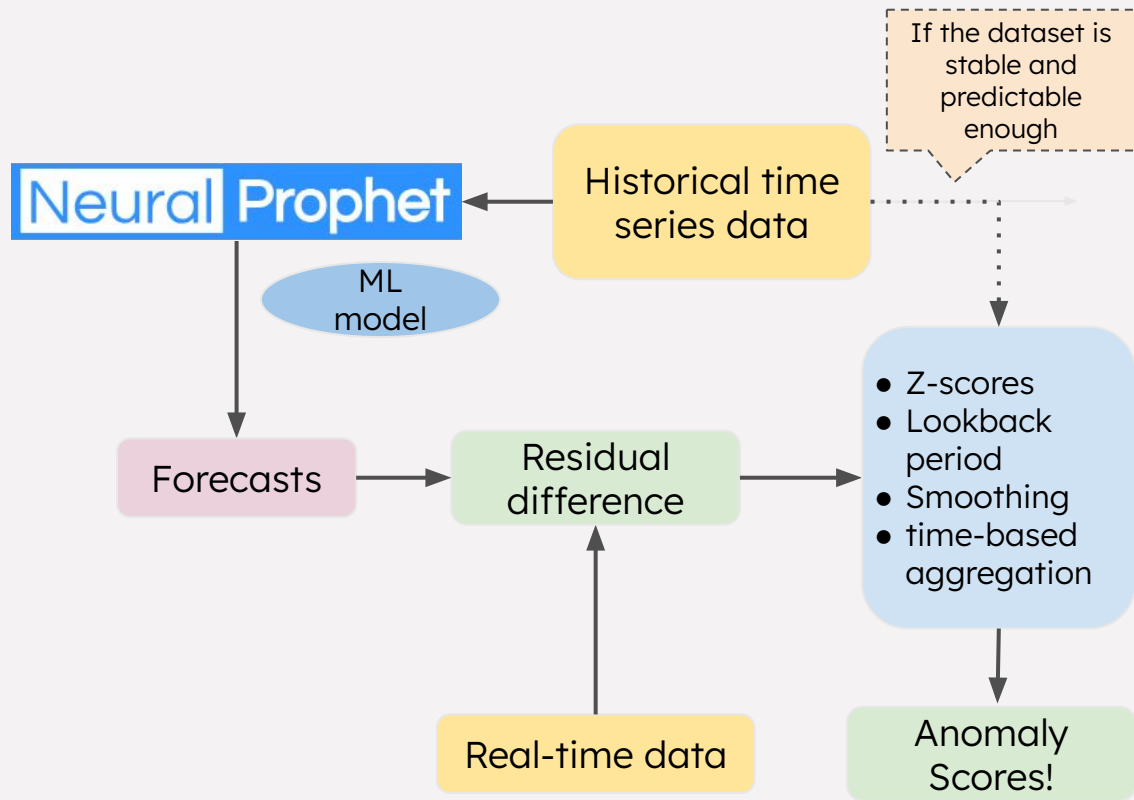
# Considerations for sparse or volatile time series

- Impute values if possible
- Aggregate into time windows
- Discrete or rolling time bucketing
- Trade-offs with latency



# Bringing it all together

- Start simple using basic z-scores and lookback period on the time series data
- If your time series is more complicated, then use forecasts to find the expected value
- Instead of modeling the time series data, model residuals
- And do all the statistical experimentation with residuals

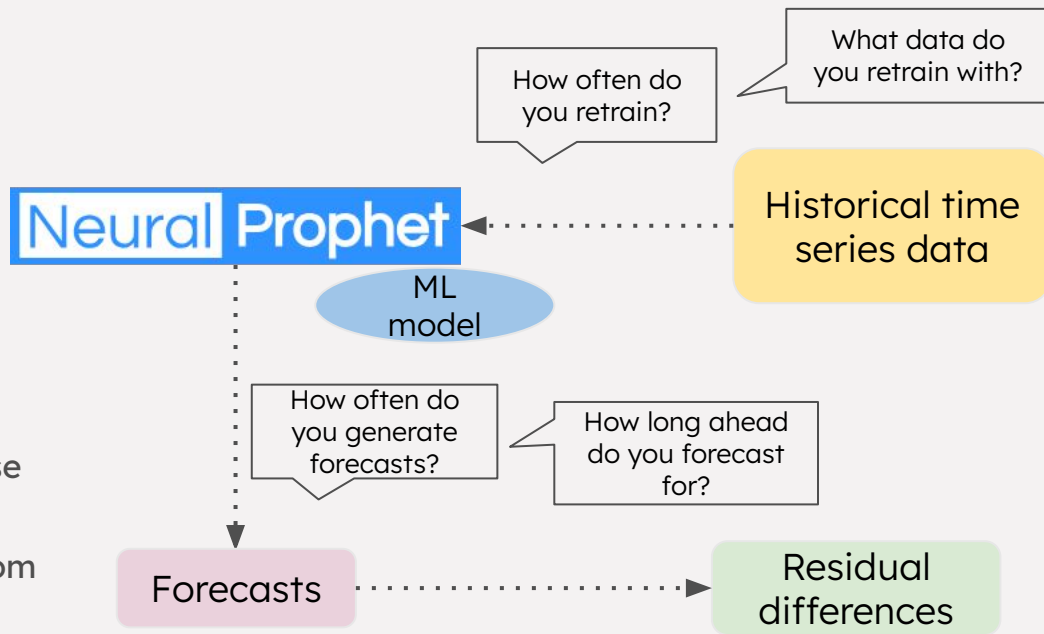


# Building a Realtime Pipeline

- Model development life cycle
- Realtime data is continuously populating your systems

## 1. Save your data from anomalies

- Anomalous data WILL creep in hence, use robust metrics
- Label and filter out anomalous period from the data to retrain ASAP
- OR remove confidently predicted anomalies

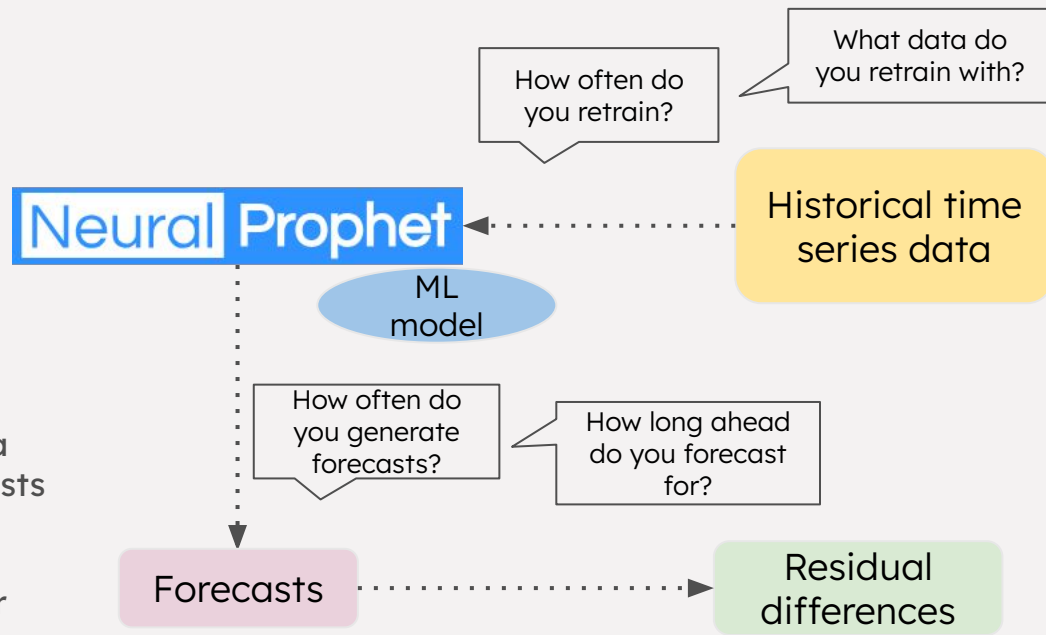




- Model development life cycle
- Realtime data is continuously populating your systems

## 2. Forecast reasonably to avoid bias

- NeuralProphet will use ALL historical data and the trained model to generate forecasts
- Forecast frequently to incorporate recent behavior
- But try to avoid using anomalous data for prediction
- OR remove confidently predicted anomalies from retraining/forecasting



# Thank you!

## Questions?

Talk contents:

<https://github.com/ShreyaKhurana/pydata-global-2024>

References:

<https://neuralprophet.com/contents.html>

<https://www.statology.org/modified-z-score/>

### Email

shreyakhurana11235@gmail.com

### LinkedIn

<https://www.linkedin.com/in/shreya-khurana/>

### GitHub

@ShreyaKhurana